



NYC-Airbnb-Recommendation-Engine-NLP / notebooks / Recommendation_Engine.ipynb

 kamalova Created using Colaboratory

History

 1 contributor

2211 lines (2211 sloc) | 199 KB

...



Airbnb Recommendation Engine for NYC through Sentiment Analysis

Table of Contents

- Business Case
- Aim of the Notebook
- What is a Recommendation Engine
- Data Understanding
- Building Recommender Engine

1. Business Case

About Airbnb: *You can host anything, anywhere, so guests can enjoy everything, everywhere.*

Nowadays the demand for short and long-term temporary accommodation is increasing thanks to easing travel conditions. This demand positively affects the number of online platforms that allow you to make reservations before traveling. **Airbnb** is one such platform, which allows travelers to make accommodation reservations based on the fact that the host leases all or part of his or her home to the traveler.

Customer reviews play an important role in the customer's decision to purchase a product or use a service. Customer preferences and opinions are affected by other customers' reviews online, on blogs or over social networking platforms

The main goal of this work is to combine both recommendation system and sentiment analysis in order to recommend the most accurate listings for users based on their preferences in **New York City**. Since both domains suffer from the lack of labeled data, to overcome that, this project detects the opinions polarity score using **NLTK VADER** (Valence Aware Dictionary and Sentiment Reasoner) Lexicon.

We'll therefore split our approaches into following sections:

- Exploring available AirBnb listings in NYC
- Measuring polarity/sentiment scores along with vader_lexicon. This polarity

measurement adapts to *pos*, *neu*, *neg*, and compound. By simply taking the compound from these values, a new feature was created on the data.

- Building a recommendation engine with Collaborative Filtering to predict sentiment score for all reviewer-listing pairs and making personalised recommendations for each user based on their ranked preferences.

2. What is a Recommendation Engine?

In general, recommendation engine consist of algorithms that can present similar elements to users. Recommended application, articles, videos, etc. It's about the user. It analyzes the user's previous habits and makes recommendations. Each item shown to the user has a ranking. This sequence is based on the recommended system and is created by examining the user's historical data. This system consists of two separate categories. **Content-Based (CB)** and **Collaborative Filtering (CF)** systems.

The CF method focuses on collecting and analyzing data on user behavior, activities, and preferences, to predict what a person will like, based on their similarity to other users.

To plot and calculate these similarities, collaborative filtering uses a matrix style formula. An advantage of collaborative filtering is that it doesn't need to analyze or understand the content (products, films, books). It simply picks items to recommend based on what they know about the user.

[more](#)

3. Aim of the Notebook

This is the last Notebook and last project section which aims to building a recommendation engine with Collaborative Filtering to predict sentiment score for all reviewer-listing pairs and make personalised recommendations for each user based on their ranked preferences.

4. Data Understanding

We will use the dataset of review_polarity which was preprocessed during the Sentiment Analysis section. Let's dive deep into the most exciting part of the project.

```
In [1]: # Import necessary Libraries
```

```
import numpy as np
import pandas as pd
pd.set_option('display.max_colwidth', None)

# Data visualization
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
# Seaborn's beautiful styling
import seaborn as sns
# to get rid of the warnings
import warnings
warnings.filterwarnings("ignore")
sns.set_style('whitegrid')
```

In [2]: `%rm -rf sample_data/`

In [3]: `# Load dataset`
`df_reviews = pd.read_csv('/content/reviews_polarity.csv')`
`df_reviews`

Out[3]:

	listing_id	id	reviewer_id	reviewer_name	comments	month	weekday	language	text_length	polarity_score
0	2595	19760	38960	Anita	i ve stayed with my friend at the midtown castle for six days and it was a lovely place to be a big spacious room with a pointy roof which really makes you feel like staying in a castle the location is perfect it is just a few steps from macy s time square	12	Thursday	en	468	0.9274

					and theatre district everything worked just perfect with the keys etc thank you so much jennifer we had a great time in new york attention it s on the floor without a lift but definetely worth it					
					we ve been staying here for about nights enjoying to be in the center of the city that never sleeps short ways to everywhere in manhattan by subway or by walk midtown					
1	2595	34320	71130	Kai-Uwe	castle is a beautiful and tastful place jennifer and tori relaxed and friendly hosts thats why we the three berliners recommand that place good to have wifi and a little kitchen too	4	Friday	en	366	0.9136
					we had a wonderful stay at iennifer s					

					very charming apartment they					
2	2595	46312	117113	Alicia	were very organized and helpful i would definitely recommend staying at the midtown castle	5	Tuesday	en	155	0.9409
					hi to everyone would say our greatest compliments to jennifer the host of midtown castle we spent in this lovely apartment in the heart of manhattan one month april and will remember this time as ours best the apartment is pretty spacious and great located the th					
3	2595	1238204	1783688	Sergey	ave right around the corner there is everything you can need during your short or long stay jennifer is very friendly vigorous and very responsible host thanks her and highly recomend this	5	Monday	en	570	0.9863

					apartment for everyone who are looking for a quiet place right in the center of the boiling midtown					
					jennifer was very friendly and helpful and her place is exactly as advertised the location is very convenient and it was a pleasure to stay at the midtown castle i definitely recommend it thanks					
4	2595	1293632	1870771	Loïc		5	Thursday	en	204	0.9542
...
70806	72265	161050979	109542482	John	vanessa was very pleasant and communication was very good	6	Friday	en	58	0.7774
70807	72265	163401732	1282541	Sofia	great location close to g train	6	Saturday	en	34	0.6249
					highly recommend cannot beat this value great location minute walk to subway and sec to bus which connects you easily and quickly to various parts of					

					manhattan and brooklyn organic as well as regular grocery stores and lots of awesome restaurants and stores near by very safe					
70808	72265	252657179	8936723	Yo	neighborhoods nice room not big but it s plenty enough and everything works well it s nice warm in the winter even though the bedroom is separated by a curtain to the kitchen because the host is mainly in the other section of the apartment you have a lot of privacy vanessa is a very friendly interesting and helpful host	4	Wednesday	en	626	0.9870
					vanessa is a great and very polite host and gives you as much privacy as you want the room can be seen in the photos and has everything you					
70809	72265	277084426	17160406	Ioannis	need the	6	Friday	en	275	0.8126

location is
amazing as
well with
plenty of bars
restaurants and
stores around
and literally
half a block
away from g
train

vanessa is a
very hospitable
and friendly
person i was
able to interact
with her a lot
the apartment
is ideally
located the
room very
convenient i
have nothing
to say except
that i had a
great time and
it was a great
experience for
me thank you
vanessa i
would
definitely go
back

70810	72265	294169497	165490874	Elsa	7	Saturday	en	287	0.9621
-------	-------	-----------	-----------	------	---	----------	----	-----	--------

70811 rows × 11 columns



```
In [4]: # Print dataframe columns
df_reviews.columns
```

```
Out[4]: Index(['listing_id', 'id', 'reviewer_id', 'reviewer_name', 'comments', 'month',
               'weekday', 'language', 'text_length', 'polarity_score',
               'sentiment_type'],
              dtype='object')
```

```
In [5]: # Drop unnecessary columns
df_reviews.drop(columns=['id', 'comments', 'month', 'weekday', 'language', 'text_length'], inplace=True)
```

```
In [6]: df_reviews
```

```
Out[6]:
```

	listing_id	reviewer_id	reviewer_name	polarity_score	sentiment_type
0	2595	38960	Anita	0.9274	Positive
1	2595	71130	Kai-Uwe	0.9136	Positive
2	2595	117113	Alicia	0.9409	Positive
3	2595	1783688	Sergey	0.9863	Positive
4	2595	1870771	Loïc	0.9542	Positive
...
70806	72265	109542482	John	0.7774	Positive
70807	72265	1282541	Sofia	0.6249	Positive
70808	72265	8936723	Yo	0.9870	Positive
70809	72265	17160406	Ioannis	0.8126	Positive
70810	72265	165490874	Elsa	0.9621	Positive

70811 rows × 5 columns

```
In [7]: df_reviews.polarity_score.describe()
```

```
Out[7]: count    70811.000000
mean         0.877807
std          0.205164
min         -0.995000
25%          0.872000
50%          0.945100
75%          0.974700
max           0.999400
Name: polarity_score, dtype: float64
```

```
In [8]: # Install surprise package
```

```
# Install surprise package
! pip install surprise
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Collecting surprise
  Downloading surprise-0.1-py2.py3-none-any.whl (1.8 kB)
Collecting scikit-surprise
  Downloading scikit-surprise-1.1.3.tar.gz (771 kB)
    |████████████████████████████████████████| 771 kB 5.0 MB/s
Requirement already satisfied: joblib>=1.0.0 in /usr/local/lib/python3.8/dist-packages (from scikit-surprise->s
urprise) (1.2.0)
Requirement already satisfied: numpy>=1.17.3 in /usr/local/lib/python3.8/dist-packages (from scikit-surprise->s
urprise) (1.21.6)
Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/python3.8/dist-packages (from scikit-surprise->su
rprise) (1.7.3)
Building wheels for collected packages: scikit-surprise
  Building wheel for scikit-surprise (setup.py) ... done
  Created wheel for scikit-surprise: filename=scikit_surprise-1.1.3-cp38-cp38-linux_x86_64.whl size=2626475 sha
256=3f9a8acdb7c934c3cc5f218132b0720143760b3169619cc39806133e536756b1
  Stored in directory: /root/.cache/pip/wheels/af/db/86/2c18183a80ba05da35bf0fb7417aac5cddb93bcb1b92fd3ea
Successfully built scikit-surprise
Installing collected packages: scikit-surprise, surprise
Successfully installed scikit-surprise-1.1.3 surprise-0.1
```

In [9]:

```
# Import an additional libraries
from surprise import SVD, Dataset, Reader, accuracy
from surprise.model_selection import cross_validate, train_test_split, GridSearchCV
```

5. Building Recommender Engine

Matrix Factorization-based Algorithm

Singular Value Decomposition(SVD) famous algorithm, as popularized by Simon Funk during the Netflix Prize.

SVD is a matrix factorization technique that is usually used to reduce the number of features of a data set by reducing space dimensions from N to K where $K < N$. For the purpose of the recommendation systems however, we are only interested in the matrix factorization part keeping same dimensionality. The matrix factorization is done on the user-item ratings matrix. From a high level, matrix factorization can be thought of as finding 2 matrices whose product is the original matrix.

Surprise package provides implementation of this algorithms.

In [10]:

```
# Rating scale is basically between -1 and 1.
```

```
reader = Reader(rating_scale=(-1,1))
```

```
In [11]: df = Dataset.load_from_df(df_reviews[['listing_id', 'reviewer_id', 'polarity_score']], reader)
```

```
In [12]: model_svd = SVD()
cv_results_svd = cross_validate(model_svd, df, cv=5)
pd.DataFrame(cv_results_svd).mean()
```

```
Out[12]: test_rmse    0.152873
test_mae     0.084486
fit_time     1.200835
test_time    0.113524
dtype: float64
```

5.1. Tuning the Algorithm Parameters

Surprise provides a GridSearchCV class analogous to GridSearchCV from scikit-learn.

With a dict of all parameters, GridSearchCV tries all the combinations of parameters and reports the best parameters for any accuracy measure.

It is used to find the best setting of parameters:

- **n_epochs** - the number of iteration *.Default is 20*
- **lr_all** - is the learning rate for all parameters, which is a parameter that decides how much the parameters are adjusted in each iteration. *Default is 0.005*
- **reg_all** - is the regularization term for all parameters, which is a penalty term added to prevent overfitting. *Default is 0.02*

As a result, regarding the majority of parameters, the default setting is the most optimal one. The improvement obtained with Grid Search is very small.

```
In [13]: # Setting dictionary parameters
param_grid = {'n_epochs': [5, 10, 20],
              'lr_all': [0.002, 0.005],
              'reg_all': [0.2, 0.4, 0.6]}

GS = GridSearchCV(SVD, param_grid, measures=['rmse', 'mae'], cv=3)
```

```
GS.fit(df)
# Best RMSE score
print(GS.best_score['rmse'])
# Combination of parameters that gave the best RMSE score
print(GS.best_params['rmse'])
```

```
0.17222471880571377
{'n_epochs': 20, 'lr_all': 0.005, 'reg_all': 0.2}
```

5.2. Analysis of Collaborative Filtering Model Results

In this part, let's examine in detail the results obtained by the SVD model that provided the best RMSE score.

```
In [14]: # Split dataset into train/test sets. Test set is made of 20% of the dataset.
train_set, test_set = train_test_split(df, test_size=0.2)

# Train the algorithm on the trainset, and predict ratings for the testset
model_svd = SVD(n_epochs=20, lr_all=0.005, reg_all=0.2)
model_svd.fit(train_set)
predictions = model_svd.test(test_set)
```

```
In [15]: print('Accuracy on test data set,', end=' ')
accuracy.rmse(predictions)
```

```
Accuracy on test data set,    RMSE: 0.1680
```

```
Out[15]: 0.16795162479164888
```

```
In [18]: df_pred = pd.DataFrame(predictions, columns=['listing_id', 'reviewer_id', 'polarity_score', 'pred_pol', 'detail

df_pred['pred_pol_round'] = df_pred['pred_pol'].round()
df_pred['abs_err'] = abs(df_pred['pred_pol'] - df_pred['polarity_score'])
df_pred.drop(['details'], axis=1, inplace=True)

df_pred.sample(5)
```

```
Out[18]:
```

listing_id	reviewer_id	polarity_score	pred_pol	pred_pol_round	abs_err
------------	-------------	----------------	----------	----------------	---------

13254	106647	147106860	0.8591	0.916416	1.0	0.057316
10902	66275	89336041	0.7650	0.861363	1.0	0.096363
8906	14434	1604432	0.9022	0.867694	1.0	0.034506
6486	64365	1357153	0.9756	0.886405	1.0	0.089195
12410	66251	168262243	0.9153	0.879173	1.0	0.036127

```
In [19]: plt.figure(figsize=(10,5))
sns.scatterplot(data=df_pred, y='polarity_score', x='pred_pol', color='#9f4e4f')
plt.title('Predicted v.s. True Polarity Scores',fontweight="bold")
plt.xlabel('Predicted Scores', fontweight="bold")
plt.ylabel('Actual Scores', fontweight="bold");
```

